

Computing on Clark and More Parallel Computing

J.D. Peiffer

2/7/2021

Section 1: Connecting to Clark with mobaXterm

1. Visit <https://missouri.app.box.com/v/rcss-mobaxterm-new> and download both files into your downloads folder.
2. Run the .exe file, accept all defaults for installation.
3. Open mobaXterm and click “Start Local Terminal” This will open a command interface for you to connect to Clark with.
4. Generate a ssh key:
 - i) In your terminal type `ssh-keygen`.
 - ii) Accept all defaults by pushing enter. Do not enter a password.
 - iii) You only need to do this once!
5. Connect to Clark

- i) In your terminal, type `ssh <username>@clark.rnet.missouri.edu`
Replace `<username>` with your pawprint.

```
11/02/2021 13:47.59 /home/mobaxterm ssh jdp6n8@clark.rnet.missouri.edu
Enter passphrase for key '/home/mobaxterm/.ssh/id_rsa':
```

- ii) Do not include brackets, for example `ssh jdp6n8@clark.rnet.missouri.edu`
- iii) When prompted for your password, enter your university password. **NOTE THAT WHAT YOU TYPE DOES NOT SHOW UP IN THE TERMINAL** You should then see something like this:

```
* *****
* Welcome to clark-r630-login-node907.cobalt.lan, CentOS, 7.9.2009
* Archive of news is available in /etc/motd-archive
* University of Missouri Division of IT Research Computing Support Services
*
* Use of this system is governed by the rules and regulations of the
* University of Missouri and the University of Missouri System and also
* by requirements of various granting agencies.
*
* Users must be familiar with and abide by the UM System acceptable use
* policy (CRR 110.005) and the UM System Data Classification System (DCL).
* http://umsystem.edu/ums/rules/collected_rules/facilities/ch110
* http://umsystem.edu/ums/is/infosec/classification
*
* Helpful Links
* News: http://doit.missouri.edu/research/research-computing-news
* Support: rcss-support@missouri.edu
* Announcement list: https://po.missouri.edu/cgi-bin/wa?A0=rcss-announce-l
* Training: http://docs.rnet.missouri.edu/training
* Documentation: http://docs.rnet.missouri.edu/
* *****
[jdp6n8@clark-r630-login-node907 ~]$
```

Section 2: Linux commands on Clark

Need to add data copy

Watch **this video** if you are not familiar with bash/linux commands. These will be useful in mobaXterm and Clark.

Question 2.1 In your own words, provide a description of what each of the following bash terminal commands do (google is your friend):

- pwd
- ls
- cd
- touch
- rm
- cat
- cp
- mv
- mkdir

Question 2.2 In your home directory, create a folder named `hw4` and in it create a file name `test.txt`. Screenshot the output of `ls` in this directory (Hint: windows users hold `windows button + shift + s`)

When you log into Clark, you are logging into part of a computer cluster called a *node*. Every time, you start out in the *login node*. Here, you can perform basic tasks like making folders and files, but more data intensive tasks, like running R or NCL, must be done on a compute node.

To request a node for this task, type this into your terminal: `srun -n 1 --mem 8G --pty /bin/bash`

Notice the node name at your command line changes!

To load the rstudio software, type `module load rstudio/1.1.456`

Next, type `rstudio`

Section 3: Load and Characterize Data

Make sure your working directory contains the file “allprcp.Rdata”.

Load the variable into your workspace with `load("allprcp.Rdata")`

Question 3.1 What are the dimensions of this dataset? What is represented on the x and y dimension? (Hint: similar to hw3).

Question 3.2 What is the beginning and end date of this time series?

Question 3.3 What is the column range of the data? (Names begin with US...) Use this information to make a new data frame called `rain`.

```
data_start_column = #Enter your answer here
data_stop_column = #Enter your answer here

rain = allprcp[, data_start_column:data_stop_column]
```

Question 3.4 How many missing data entries are there in the entirety of this dataset? (Hint: `sum()` and `is.na()`)

That's a lot of missing data! Let us try and identify a suitable date to begin our analysis. If we start at 1980 instead of 1970, there may be less missing data! Remember what we did in hw3.

Section 4: Quantify NAs in Serial

This section will use the same code as hw3, but since the clark cluster has many cores, it allows us to highly parallelize this task!

First, let us run the code to quantify NAs with only one processor! This may take a **while**.

```
addNA = function(data_with_some_nas, starting_row) {
  nas = is.na(data_with_some_nas[starting_row:dim(data_with_some_nas)[1], ])
  return(apply(nas, 2, sum))
}

NAs_in_rain = rain * 0
system.time({
  for (i in 1:dim(rain)[1]) {
    NAs_in_rain[i, ] = addNA(rain, i)
  }
})
```

Question 4.1 How long did this take to run on clark? (Do not try on your personal computer).

Section 5: Quantify NAs in Parallel

Exit your rstudio session and type `exit` in your command prompt. This will exit your compute node. To run parallel jobs, we need to request more cores!

Type `srunch -n 8 --mem 8G --pty /bin/bash`. Load rstudio and your data as done previously. Now we will install the parallel package and parallelize our code.

```
install.packages("doParallel")
library(doParallel)

registerDoParallel()
getDoParWorkers()
```

Question 5.1 What is the output of `getDoParWorkers()`? What does this represent?

Now, run this task using two cores!

```
registerDoParallel(cores=2)
system.time({
  NAs_in_rain_par = foreach(i = 1:dim(allprcp)[1],
    .export = c("addNA", "rain"),
    .combine = "rbind") %dopar% {
    addNA(rain, i)
  }
})
```

Question 5.2 What is the elapsed time of this code? Is it faster than our earlier serial code? Why is this?
Question 5.2 Run the same code with 4 cores, 6 cores, and 8 cores. Record and plot the output with cores on the x axis and time on the y. Use any plotting method you like. **Question 5.3** Describe your graph.